

EcoMart Relational Database Design

A. Scenario 2 - EcoMart

A1: Business Problem: The EcoMart company faces challenges in managing and analyzing large-scale sales data across various regions, countries, and product categories. The current manual process is time-consuming, prone to errors, and lacks the scalability to accommodate growing data volumes. A database solution is essential to streamline data storage, enable efficient querying, and generate actionable insights.

A2: Proposed Data Structure: The proposed data structure includes three core tables—Categories, Products, and Transactions—each designed to store specific attributes. These tables are interlinked through primary and foreign keys, ensuring data consistency and enabling detailed analysis.

A3: Justification for Database Solution: A relational database offers scalability, data integrity, and efficient query performance. It eliminates redundancy through normalization, supports large-scale data processing, and provides a robust platform for advanced analytics. By using PostgreSQL, EcoMart can implement a cost-effective, reliable, and scalable solution.

A4. Usage of Business Data: The business data will be used to analyze sales trends by region and product category, monitor stock levels, evaluate revenue performance, and optimize order fulfillment. It will also support decision-making through real-time reporting and analytics.

B. Logical Data Model

The logical data model represents the structure of the database, including entities, attributes, and relationships. It ensures data is organized efficiently, adheres to normalization standards, and supports seamless querying and analysis.

Entities in the model include Categories, Products, and Transactions. Each entity contains attributes tailored to its role, such as `category_name`, `price`, and `order_date`. Relationships between tables ensure data integration, such as the one-to-many relationship between Categories and Products.

Entities and Attributes:

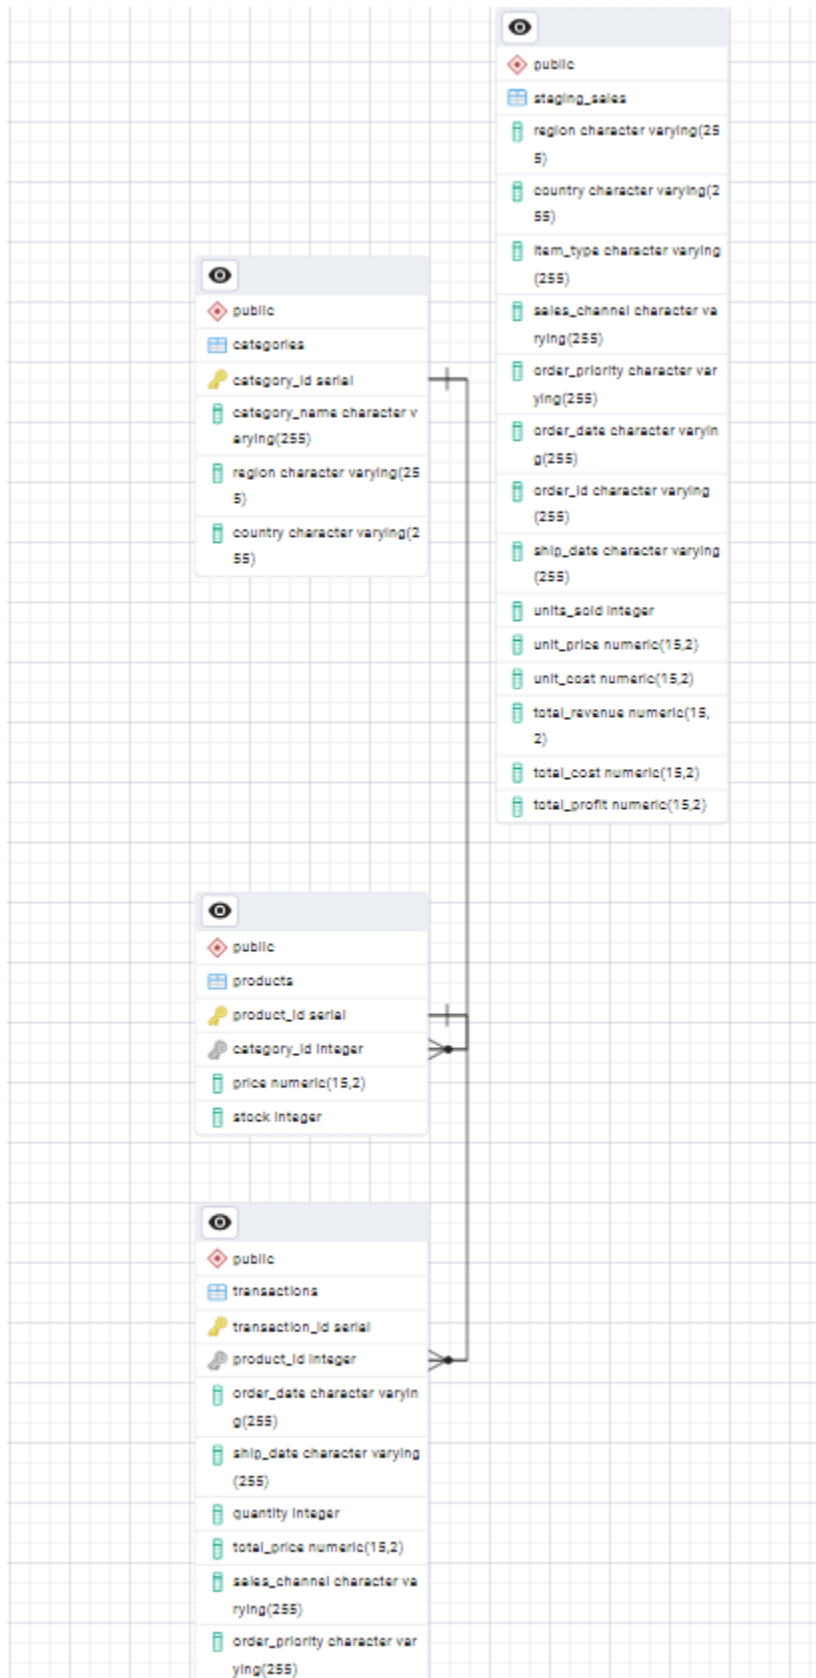
- Categories: `category_id`, `category_name`, `region`, and `country`.
- Products: `product_id`, `category_id`, `price`, and `stock`.
- Transactions: `transaction_id`, `product_id`, `order_date`, `ship_date`, `quantity`, and `total_price`.

Relationships:

- Categories → Products: One category can have multiple products.
- Products → Transactions: One product can have multiple transactions.

Normalization:

The design adheres to 3NF, ensuring no redundant data and maintaining data integrity.



C. Database Objects and Storage

The database consists of core objects like tables, indexes, and relationships designed to store and integrate data effectively. Tables represent entities, with attributes optimized for storage efficiency. For instance, “**NUMERIC**” is used for monetary values, while “**VARCHAR**” handles textual data. Indexes are automatically created for primary and foreign keys to optimize query performance.

D. Scalability

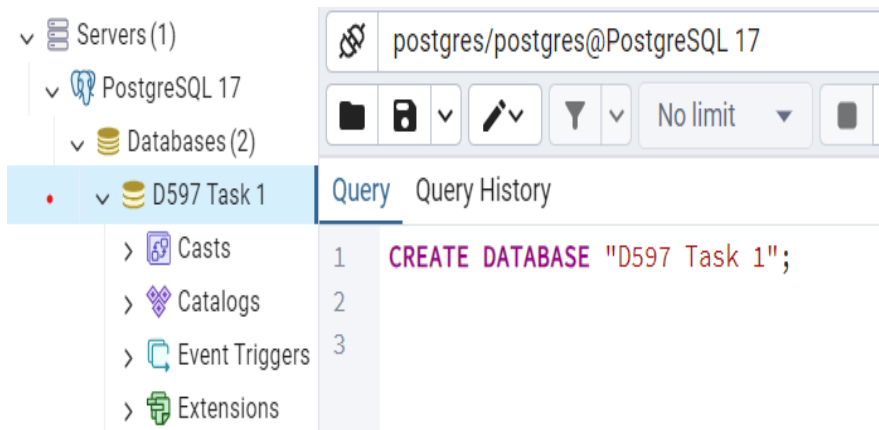
The database design addresses scalability concerns to handle increasing data volumes and user demands. Key strategies include indexing for optimized query performance, particularly on fields like “**order_date**” and “**region**”. Normalization minimizes redundancy, enhancing data consistency and storage efficiency. This makes data easier to scale than it was previously, and no need to go through thousands of rows of data.

E. Privacy and Security

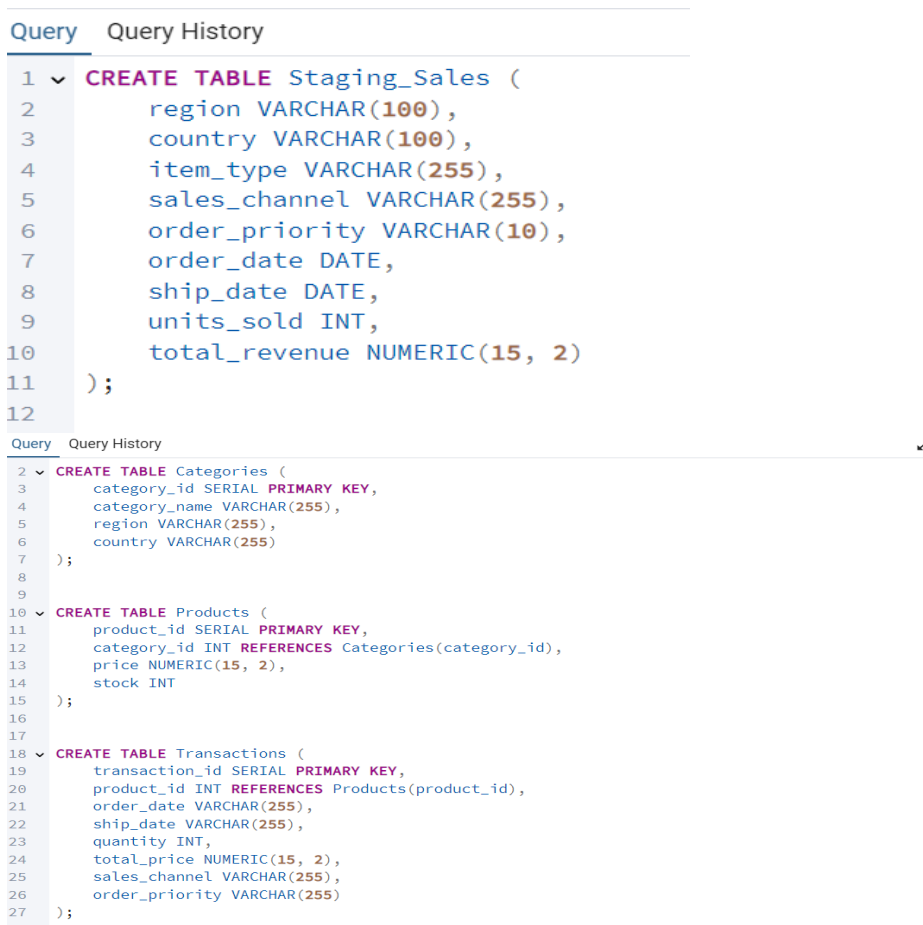
Strong privacy and security measures are built into the database design to ensure compliance with regulations and protect sensitive information. Encryption safeguards data both at rest and in process, preventing unauthorized access. The company can also require audit logging and 2FA as an added layer of security precautions. Additionally, input validation is implemented to protect against SQL attacks, which will assist in maintaining the database’s integrity.

F. Implementation Through pgAdmin

F1: Creation of Database:



F1: Creation of Tables:



F2. Data import from “Staging Table” to respective tables:

Query Query History

```
1 INSERT INTO Categories (category_name, region, country)
2 SELECT DISTINCT item_type, region, country FROM Staging_Sales;
3
4 SELECT * FROM Categories LIMIT 10;
5
6
7
```

Data Output Messages Notifications

	category_id [PK] integer	category_name character varying (255)	region character varying (255)	country character varying (255)
1	1	Household	Australia and Oceania	Papua New Guinea
2	2	Snacks	Australia and Oceania	East Timor
3	3	Vegetables	Sub-Saharan Africa	South Sudan
4	4	Beverages	Asia	Maldives
5	5	Cosmetics	Australia and Oceania	Australia
6	6	Personal Care	Sub-Saharan Africa	Sao Tome and Principe
7	7	Clothes	Asia	Singapore

Total rows: 10 Query complete 00:00:00.175

Query Query History

```
1 INSERT INTO Products (category_id, price, stock)
2 SELECT c.category_id,
3        AVG(ss.total_revenue / NULLIF(ss.units_sold, 0)),
4        SUM(ss.units_sold)
5 FROM Staging_Sales ss
6 JOIN Categories c ON ss.item_type = c.category_name
7 GROUP BY c.category_id;
8
9 SELECT * FROM Products LIMIT 10;
```

Data Output Messages Notifications

	product_id [PK] integer	category_id integer	price numeric (15,2)	stock integer
1	1	1798	205.70	42254418
2	2	1489	81.73	41517766
3	3	1269	47.45	41514213
4	4	652	47.45	41514213
5	5	273	47.45	41514213
6	6	1560	255.28	41911620
7	7	51	154.06	41254836

Total rows: 10 Query complete 00:00:00.027

Query Query History

```
1  ✓ INSERT INTO Transactions (product_id, order_date, ship_date, quantity, total_price, sales_channel, order_priority)
2  SELECT p.product_id,
3         ss.order_date,
4         ss.ship_date,
5         ss.units_sold,
6         ss.total_revenue,
7         ss.sales_channel,
8         ss.order_priority
9  FROM Staging_Sales ss
10 JOIN Products p ON ss.item_type = (SELECT category_name FROM Categories WHERE category_id = p.category_id);
11
12 SELECT * FROM Transactions LIMIT 10;
13
```

[Data Output](#) [Messages](#) [Notifications](#)

Showing rows: 1 to 10								
Page No: 1								
	transaction_id [PK] integer	product_id integer	order_date character varying (255)	ship_date character varying (255)	quantity integer	total_price numeric (15,2)	sales_channel character varying (255)	order_priority character varying (255)
1	1	632	12/27/2012	1/30/2013	2910	742864.80	Online	M
2	2	632	4/26/2010	5/25/2010	2302	587654.56	Online	H
3	3	632	5/5/2011	5/13/2011	6574	1678210.72	Online	C
4	4	632	7/10/2012	7/22/2012	8669	2213022.32	Online	M
5	5	632	11/24/2015	12/1/2015	1368	349223.04	Offline	M
6	6	632	10/31/2013	11/7/2013	4209	1074473.52	Online	C

F3: Query Examples:

F3: Query 1 - Total Revenue by Region:

A query for Total Revenue by Region helps analyze which areas are performing the best or need improvement, so decisions about where to focus resources can be made. It's useful for planning strategies, budgeting, and forecasting future performance, while also identifying potential growth opportunities in certain regions. This type of query can also support compliance, reporting needs, and provide insights for comparing performance against competitors, helping the business stay on track with its goals.

Query

Query History

1

▼

SELECT region, **SUM**(total_price) **AS** total_revenue

2

FROM Transactions t

3

JOIN Products p **ON** t.product_id = p.product_id

4

JOIN Categories c **ON** p.category_id = c.category_id


5


GROUP BY region;


Data Output


Messages


Notifications























SQL

	region character varying (255)	total_revenue numeric
1	Asia	14429520691172.28
2	Australia and Oceania	8016400383984.60
3	Central America and the Caribbean	10688533845312.80
4	Europe	25652481228750.72
5	Middle East and North Africa	12291813922109.72
6	North America	2137706769062.56
7	Sub-Saharan Africa	25652481228750.72

Before optimization:

Query

Query History

1

2

3

4

5

6

7

```

EXPLAIN ANALYZE
SELECT region, SUM(total_price) AS total_revenue
FROM Transactions t
JOIN Products p ON t.product_id = p.product_id
JOIN Categories c ON p.category_id = c.category_id
GROUP BY region;

```

Data Output

Messages

Notifications

+

📄

▼

📋

▼

🗑️

🔄

📶

SQL

Showing rows: 1 to 28

✎

QUERY PLAN

text

🔒

21

22

23

24

25

26

27

28

-> Hash (cost=109.60..109.60 rows=6660 width=8) (actual time=1.434..1.435 rows=6660 loops=3)

Buckets: 8192 Batches: 1 Memory Usage: 325kB

-> Seq Scan on products p (cost=0.00..109.60 rows=6660 width=8) (actual time=0.061..0.749 rows=6660 loops=3)

-> Hash (cost=84.40..84.40 rows=4440 width=20) (actual time=1.124..1.125 rows=4440 loops=3)

Buckets: 8192 Batches: 1 Memory Usage: 301kB

-> Seq Scan on categories c (cost=0.00..84.40 rows=4440 width=20) (actual time=0.080..0.568 rows=4440 loops=3)

Planning Time: 4.095 ms

Execution Time: 31408.408 ms

Total rows: 28

Query complete 00:00:32.395

After optimization:

Query Query History

```

1  EXPLAIN ANALYZE
2  WITH pre_aggregated AS (
3      SELECT product_id, SUM(total_price) AS total_revenue
4      FROM Transactions
5      GROUP BY product_id
6  )
7  SELECT c.region, SUM(pa.total_revenue) AS total_revenue
8  FROM pre_aggregated pa
9  JOIN Products p ON pa.product_id = p.product_id
10 JOIN Categories c ON p.category_id = c.category_id
11 GROUP BY c.region;
12
13

```

Data Output Messages Notifications



Showing rows: 1 to 34

QUERY PLAN		text	
27	Worker 0: Batches: 1 Memory Usage: 1489kB		
28	Worker 1: Batches: 1 Memory Usage: 1489kB		
29	-> Parallel Seq Scan on transactions (cost=0.00..1070773.33 rows=30833333 width=12) (actual time=0.504.....		
30	-> Hash (cost=84.40..84.40 rows=4440 width=20) (actual time=0.714..0.714 rows=4440 loops=1)		
31	Buckets: 8192 Batches: 1 Memory Usage: 301kB		
32	-> Seq Scan on categories c (cost=0.00..84.40 rows=4440 width=20) (actual time=0.010..0.290 rows=4440 loops=1)		
33	Planning Time: 0.221 ms		
34	Execution Time: 13983.656 ms		

Runtime speed changed from 31408.4 ms to 13983.6 ms

F3: Query 2 - Top 5 Selling Products

A query for the Top 5 Selling Products helps identify which products are performing the best, giving insight into what customers prefer. This information is useful for making decisions about inventory, marketing, and production, ensuring resources are focused on the most profitable items. It also helps with forecasting future demand, planning promotions, and understanding trends. By knowing your top-performing products, you can drive more sales, increase profitability, and align strategies with what works best for the business.

Query

Query History

```

1  SELECT product_id, SUM(quantity) AS total_sold
2  FROM Transactions
3  GROUP BY product_id
4  ORDER BY total_sold DESC
5  LIMIT 5;

```

Data Output

Messages

Notifications

+

📄

▼

📋

▼

🗑️

🗄️

⬇️

📈

SQL









	product_id integer	total_sold bigint
1	75	84586660
2	78	84586660
3	42	84586660
4	60	84586660
5	91	84586660

Before Optimization:


Query Query History


```
1  EXPLAIN ANALYZE
2  SELECT product_id, SUM(quantity) AS total_sold
3  FROM Transactions
4  GROUP BY product_id
5  ORDER BY total_sold DESC
6  LIMIT 5;
```

Data Output Messages Notifications

SQL

Showing rows: 1 to 22 

<div style="display: flex; justify-content: space-between; align-items: center;"> QUERY PLAN  </div>	
14	text
15	worker 1: Sort Method: quicksort Memory: 197kB
16	-> Partial HashAggregate (cost=1224940.00..1225002.52 rows=6252 width=12) (actual time=11175.922..11177.618 rows=31...
17	Group Key: product_id
18	Batches: 1 Memory Usage: 465kB
19	Worker 0: Batches: 1 Memory Usage: 465kB
20	Worker 1: Batches: 1 Memory Usage: 465kB
21	-> Parallel Seq Scan on transactions (cost=0.00..1070773.33 rows=30833333 width=8) (actual time=0.447..4.667 row...
22	Planning Time: 0.624 ms
23	Execution Time: 11224.192 ms

Total rows: 22
Query complete 00:00:11.599

After optimization:

Query

Query History

1

2

3

4

5

6

EXPLAIN ANALYZE

SELECT product_id, SUM(quantity) AS total_sold

FROM Transactions

GROUP BY product_id

ORDER BY total_sold DESC

LIMIT 5;

Data Output

Messages

Notifications

+

SQL

Showing rows: 1 to 22

14

15

16

17

18

19

20

21

22

QUERY PLAN

text

worker 1: Sort Method: quicksort Memory: 197kB

-> Partial HashAggregate (cost=1224940.00..1225002.52 rows=6252 width=12) (actual time=11158.442..11159.266 rows=31...

Group Key: product_id

Batches: 1 Memory Usage: 465kB

Worker 0: Batches: 1 Memory Usage: 465kB

Worker 1: Batches: 1 Memory Usage: 465kB

-> Parallel Seq Scan on transactions (cost=0.00..1070773.33 rows=30833333 width=8) (actual time=1.245..4673.520 row...

Planning Time: 3.463 ms

Execution Time: 11218.343 ms

Runtime speed changed from 11224.1 ms to 11218.3 ms

F3: Query 3 - Analysis of Monthly Revenues

A query for the Analysis of Monthly Revenues helps track how the business is performing over time and spot trends or patterns in revenue. This information is useful for identifying peak sales periods, managing cash flow, and setting realistic goals for growth. It also helps in comparing month-to-month performance, forecasting future revenue, and planning strategies to address any dips or maximize high-performing months. Understanding monthly revenue trends is key to staying on top of the business's financial health and making informed decisions.

[Query](#)
[Query History](#)

1
2
3
4
5
6
7
8
9
10
11
12
13

SELECT
DATE_TRUNC('month', CAST(order_date AS DATE)) AS Month,
SUM(total_price) AS Monthly_Revenue
FROM Transactions
GROUP BY DATE_TRUNC('month', CAST(order_date AS DATE))
ORDER BY Month;

[Data Output](#)
[Messages](#)
[Notifications](#)

SQL

	month timestamp with time zone	monthly_revenue numeric
1	2010-01-01 00:00:00-08	1059006757856.80
2	2010-02-01 00:00:00-08	1000411955904.80
3	2010-03-01 00:00:00-08	1025783010883.00
4	2010-04-01 00:00:00-07	1105541270873.80
5	2010-05-01 00:00:00-07	1038647792105.60
6	2010-06-01 00:00:00-07	1051420657025.40
7	2010-07-01 00:00:00-07	1082254221004.40
8	2010-08-01 00:00:00-07	1189199183456.60
9	2010-09-01 00:00:00-07	1162550268306.60

Total rows: 91
Query complete 00:01:46.561

Before optimization:

Query

Query History

1

2

3

4

5

6

7

8

9

10

11

12

13

14

EXPLAIN ANALYZE

SELECT

DATE_TRUNC('month', CAST(order_date AS DATE)) AS Month,

SUM(total_price) AS Monthly_Revenue

FROM Transactions

GROUP BY DATE_TRUNC('month', CAST(order_date AS DATE))

ORDER BY Month;

Data Output

Messages

Notifications

SQL

Showing rows: 1 to 18

Page 1

QUERY PLAN

text

11

12

13

14

15

16

17

18

-> Partial HashAggregate (cost=1533273.33..1533335.54 rows=2765 width=40) (actual time=120823.803..120823.879 rows=91 loops=3)

Group Key: date_trunc('month'::text, ((order_date)::date)::timestamp with time zone)

Batches: 1 Memory Usage: 177kB

Worker 0: Batches: 1 Memory Usage: 177kB

Worker 1: Batches: 1 Memory Usage: 177kB

-> Parallel Seq Scan on transactions (cost=0.00..1379106.66 rows=30833333 width=16) (actual time=1.366..106599.275 rows=2466...

Planning Time: 0.104 ms

Execution Time: 120867.109 ms

Total rows: 18

Query complete 00:02:01.048

After optimization:

Query

Query History

1

EXPLAIN ANALYZE

2

SELECT

3

DATE_TRUNC('month', CAST(order_date AS DATE)) AS Month,

4

SUM(total_price) AS Monthly_Revenue

5

FROM Transactions

6

GROUP BY DATE_TRUNC('month', CAST(order_date AS DATE))

7

ORDER BY Month;

Data Output

Messages

Notifications

SQL

Showing rows: 1 to 18

Page No:

QUERY PLAN

text

11

-> Partial HashAggregate (cost=1533273.33..1533335.54 rows=2765 width=40) (actual time=117559.831..117559.916 rows=91 loops=3)

12

Group Key: date_trunc('month'::text, ((order_date)::date)::timestamp with time zone)

13

Batches: 1 Memory Usage: 177kB

14

Worker 0: Batches: 1 Memory Usage: 177kB

15

Worker 1: Batches: 1 Memory Usage: 177kB

16

-> Parallel Seq Scan on transactions (cost=0.00..1379106.66 rows=3083333 width=16) (actual time=0.539..104079.379 rows=2466...

17

Planning Time: 0.388 ms

18

Execution Time: 117618.216 ms

Runtime speed changed from 120867.1 ms to 117618.2 ms

F4: Optimization Techniques:

[Query](#) [Query History](#)

1

CREATE INDEX idx_transactions_order_date ON Transactions(order_date);

2

CREATE INDEX idx_product_id ON Products (product_id);

3

CREATE INDEX idx_region ON Transactions (region);

Data Output

[Messages](#)

Notifications

CREATE INDEX

Query returned successfully in 5 min 49 secs.

G: Panopto Video

H: Sources

No sources, web or otherwise, were used in the creation of the project and its follow up report.